

# Chapter : Process Synchronization

# Background

- ❑ **Co-Operating Process:** that can affect or be affected by other processes executing in system
- ❑ Concurrent access to shared data may result in data inconsistency
- ❑ **Process Synchronization:** Ensures coordination among processes and maintains Data Consistency

## ❑ Process P1

1.  $X=5$
2.  $X=5+2$

## Process P2

1.  $\text{read}(x);$
2.  $x=x+5;$

3.  $\text{Printf}(x);$



# **Producer Consumer Problem**

**There can be two situations:**

- 1. Producer Produces Items at Fastest Rate Than Consumer Consumes**
- 2. Producer Produces Items at Lowest Rate Than Consumer Consumes**



# Producer Consumer Problem

**Producer Produces Items at Fastest Rate Than  
Consumer Consumes:**

If Producer produces items at fastest rate than  
Consumer consumes Then Some items will be  
lost

Eg. Computer → Producer


Printer → Consumer

# Producer Consumer Problem

## Solution:

To avoid mismatch of items Produced or Consumed →  
Take Buffer

Idea is: Instead of sending items from Producer to  
Consumer directly → Store items into buffer



# Producer Consumer Problem

Buffer Can be:

## 1. Unbounded Buffer:

1. No buffer size limit
2. Any no. of items can be stored
3. Producer can produce on any rate, there will always be space in buffer

## 2. Bounded Buffer:

1. Limited buffer size

# Producer Consumer Problem

Bounded Buffer:

If rate of Production  $>$  rate of Consumption:

Some items will be unconsumed in buffer

If rate of Production  $<$  rate of Consumption:

At some time buffer will be empty

# Producer

```
while (true) {
```

```
    /* produce an item and put in  
    nextProduced */
```

```
    while (count == BUFFER_SIZE)  
        ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;
```

```
}
```



# Consumer

```
while (true) {  
    while (count == 0)        // buffer empty  
    {  
        ; // do nothing  
    }  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count- -;  
  
    /* consume the item in nextConsumed  
    */  
}
```

# Race Condition

- When multiple processes access and manipulate the same data at the same time, they may enter into a race condition.
- **Race Condition: When output of the process is dependent on the sequence of other processes.**
- Race Condition occurs when processes share same data

## □ Process P1

1. reads  $i=10$
2.  $i=i+1 =11$

## Process P2

1. P2 reads  $i=11$  from memory
  2.  $i=i+1 = 12$
  3. Stores 12 in memory
3. Stores  $i=11$  in memory

# Critical Section Problem

- Section of code or set of operations, in which process may be changing shared variables, updating common data.
- **A process is in the critical section if it executes code that manipulate shared data and resources.**
- Each process should seek permission to enter its critical section → **Entry Section**
- **Exit Section**
- **Remainder section:** Contains remaining code

# Structure of a process

**Repeat**

**{**

**Locks are set here**

**// Entry Section**

**Critical Section** ( a section of code  
where processes work with shared  
data)

**Critical Section**

**Locks are released here**

**// Exit Section**

**Remainder Section**

**} until false.**

# Solution to: Critical Section

## 1. Mutual Exclusion:

It states that **if one process is executing in its critical section**, then no other process can execute in its critical section.

## 2. Bounded Wait:

It states that if a process makes a request to enter its critical section and before that request is granted, **there is a limit on number of times other processes are allowed to enter that critical section**.

## 3. Progress:

It states that **process cannot stop other process from entering their critical sections, if it is not executing in its CS**.

# Peterson's Solution

- Two process solution (Software-based)
- The two processes share two variables:
  - int **turn**;
  - Boolean **flag[2]**
- The variable **turn** indicates whose turn it is to enter the critical section.
- The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i]=true** implies that process **P<sub>i</sub>** is ready!

- Only **2 processes**,
- General structure of process  $P_i$  (other process  $P_j$ )  
**do** {  
    **entry** section  
    critical section  
    **exit** section  
    reminder section  
} **while (1);**

# Algorithm for Process $P_i$

do {

flag[i] = TRUE;

turn = j;

while ( flag[j] && turn == j);

CRITICAL SECTION

flag[i] = FALSE;

REMAINDER SECTION



do {

*acquire lock*

critical section

*release lock*

remainder section

} while (TRUE);

# Critical Section Problem solution

- Now prove that this solution is correct.  
We need to show that:
  - Mutual exclusion is preserved.
  - The progress requirement is satisfied.
  - The bounded-waiting requirement is met.

# Synchronization Hardware

## Hardware Solution to C.S.

- ❑ Many systems provide hardware support for critical section code
- ❑ Uni-processors – could disable interrupts
  - ❑ Currently running code would execute without preemption
- ❑ Modern machines provide special atomic hardware instructions
  - ▶ **Atomic = non-interruptable**
  - ❑ Either test memory word and set value
  - ❑ Or swap contents of two memory words

## 1. Interrupt Disabling

**1. Process leaves control of CPU when it is interrupted.**

**2. Solution is:**

- 1. To have each process disable all interrupts just after entering to the critical section.**
- 2. Re-enable interrupts after leaving critical section**

## □ Interrupt Disabling

**Repeat**

**Disable interrupts**

**C.S**

**Enable interrupts**

**Remainder section**

# Synchronization Hardware

- Having the support of some simple hardware instructions, the CS problem can be solved very easily and efficiently.
- The CS problem occurs because the modification of a shared variable of a process may be interrupted.
- Two common hardware instructions that execute **atomically**
  - **Test-and-Set**
  - **Swap**

# Synchronization Hardware

- Test and modify the content of a word atomically.

```
boolean TestAndSet(boolean &target) {  
    boolean rv = target;  
    target = true;  
  
    return rv;  
}
```

# Mutual Exclusion with Test-and-Set

- Shared data:  
**boolean lock = false;**
- Process  $P_i$   
**do {**  
    **while (TestAndSet(lock)) ;**  
    critical section  
    **lock = false;**  
    remainder section  
**}**



# Synchronization Hardware

- Atomically swap two variables.

```
void Swap(boolean &a, boolean  
&b) {  
    boolean temp = a;  
    a = b;  
    b = temp;  
}
```

# Mutual Exclusion with Swap



- Shared data (initialized to **false**):  
boolean lock; /\***global** variable
- Process  $P_i$   
do {  
    **key = true;**  
    while (key == true)  
        Swap(lock, key);  
    critical section  
    **lock = false;**  
    remainder section  
}

# Semaphore

- Synchronization tool that maintains concurrency using variables
- Semaphore S is a integer variable which can take positive values including 0. It is accessed from 2 operations only.

## Operations On Semaphore:

- `wait()` and `signal()`

1. Wait Operation is also known as P() which means to test
2. Signal Operation is also known as V() which means to increment

- Entry to C.S. is controlled by `wait()`
- Exit from C.S. is signaled by `signal()`

- Can only be accessed via two indivisible (atomic) operations and  $S=1$
- For critical Section problem semaphore value is always 1
- $P(S)$  and  $V(S)$ :

```
wait (S) {  
    while (S <= 0)  
    do skip ; // no-op}  
    S- -;
```

**C.S**

```
signal (S) {  
    S++;    }
```

# Semaphore as General Synchronization Tool

## Types of Semaphores:

- **Counting** semaphore – when integer value can be any non-negative value
- **Binary** semaphore – integer value can range only between 0 and 1
  - Also known as **mutex locks**

# Semaphore and Busy Waiting

## Disadvantage of Semaphore: Busy Waiting

- ❑ When a process is in C.S. and any other process that wants to enter C.S. loops continuously in entry section.
- ❑ Wastes CPU cycles
- ❑ Semaphore that implements busy waiting is known as: **Spin Lock**

# Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated **waiting queue**. Each entry in a waiting queue has two data items:
  - value (of type integer)
  - pointer to next record in the list

# Semaphore Implementation with no Busy waiting

Instead of waiting, a process blocks itself.

- Two operations:
  - **block** – place the process invoking the operation on the waiting queue.
  - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue.



# Semaphore Implementation with no Busy waiting

□ Implementation of wait:

$S=1$

```
wait (S){  
    value--;  
    if (value < 0) {  
        add process P to waiting queue  
        block();  
    }  
C.S  
}
```

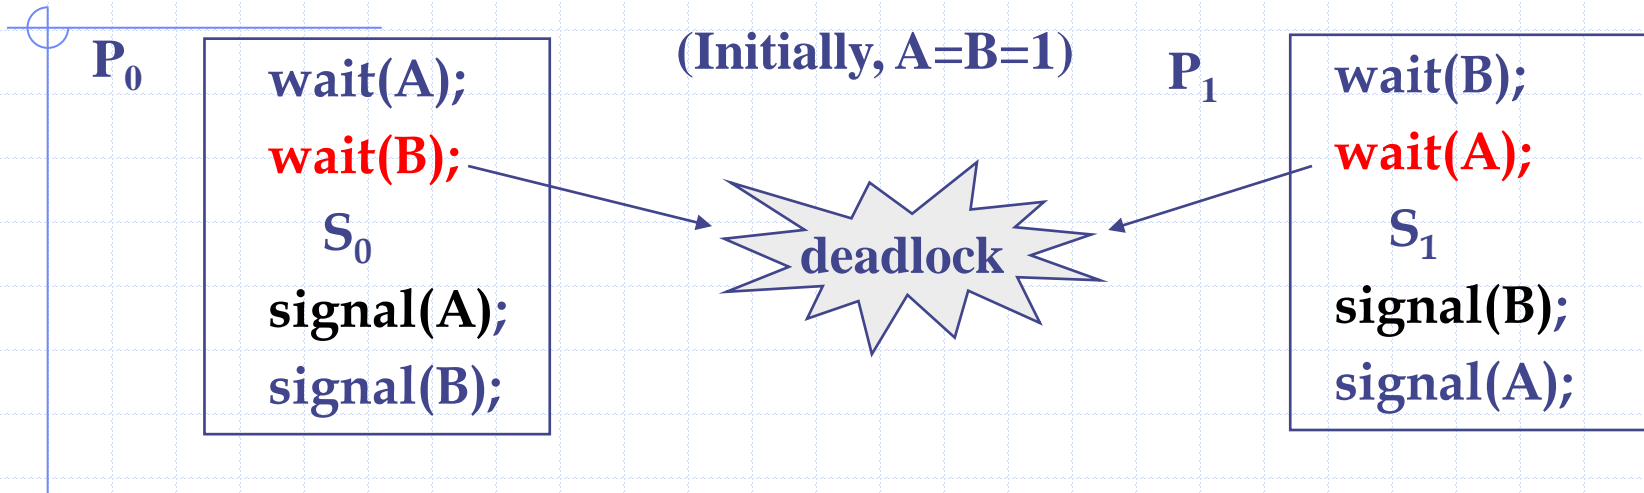
# Semaphore Implementation with no Busy waiting

- Implementation of signal:

```
Signal (S){  
    value++;          //“value” has -ve value i.e -1 here  
    if (value <= 0) {  
        remove a process P from the waiting queue  
        wakeup(P); }  
}
```

# Deadlocks and Starvation

- ◆ The using of semaphores may cause **deadlocks**



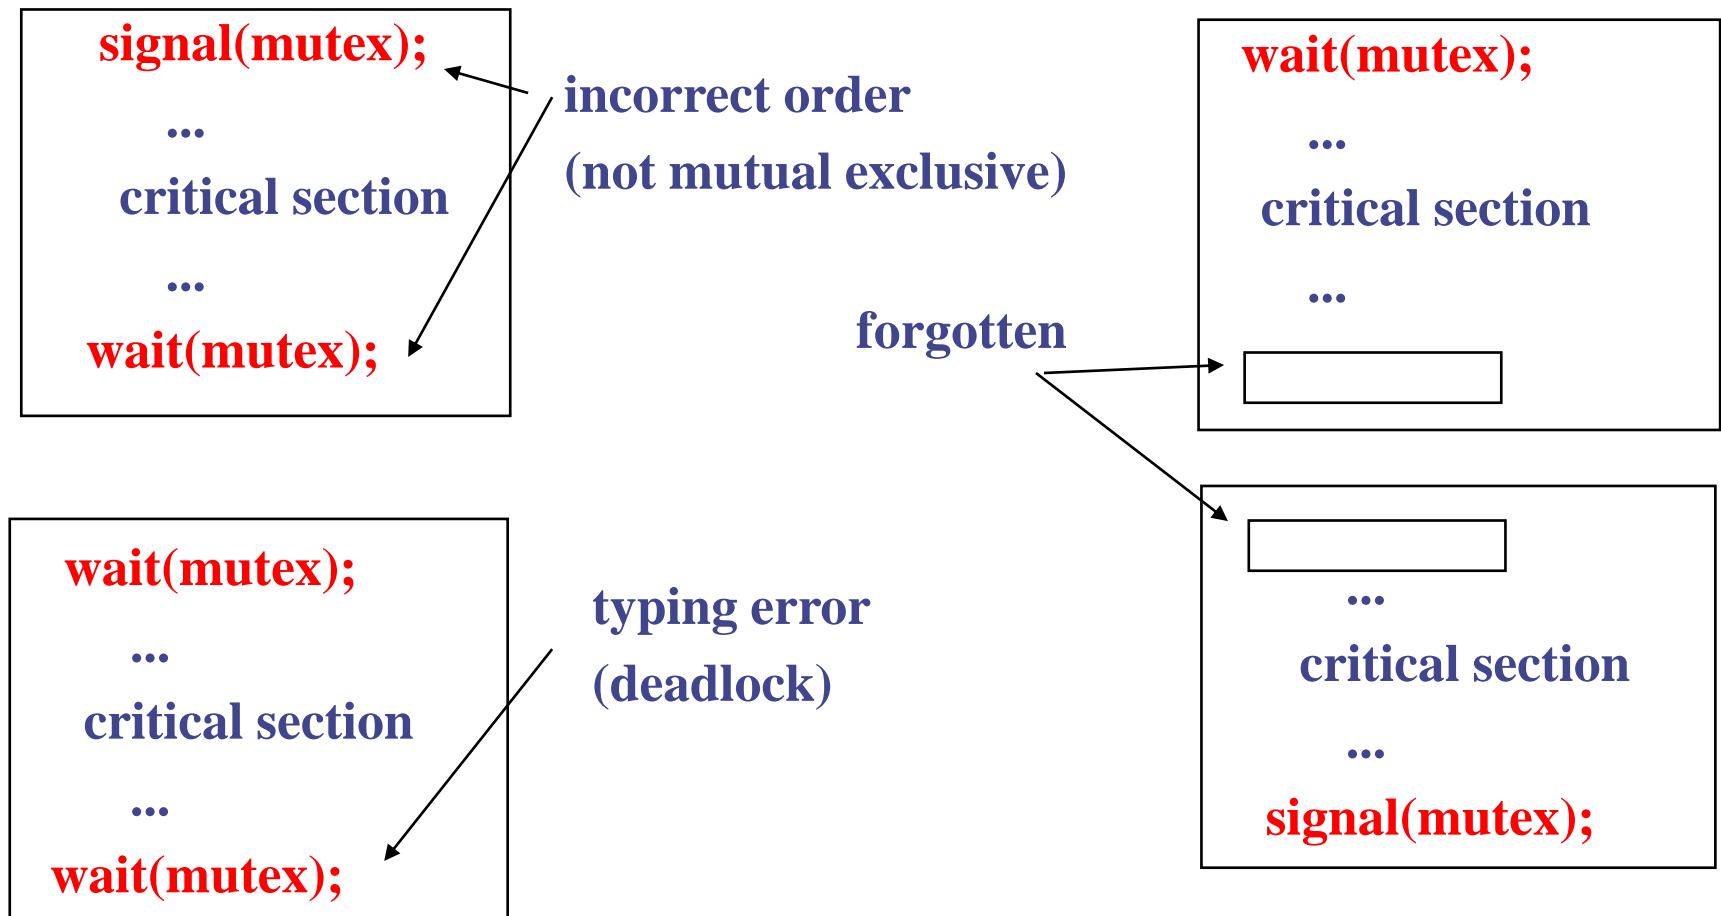
- ◆ **Starvation** – **indefinite blocking**. A process may never be removed from the semaphore queue in which it is suspended

For example: waiting queues are implemented in LIFO order.

# Drawbacks of Semaphores



- Semaphores provide a convenient and effective mechanism for process synchronization.
- However, **incorrect use** may result in timing errors.



# Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

# Readers-Writers Problem

- A data set is shared among a number of concurrent processes
  - Readers – only read the data set; they do **not** perform any updates
  - Writers – can both read and write.
- Problem – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.
- Shared Data
  - Data set
  - For Readers: Semaphore **mutex** initialized to 1.
  - For Writers: Semaphore **wrt** initialized to 1.
  - Integer **readcount** initialized to 0.

# Readers-Writers Problem (Cont.)

- The structure of a writer process

```
while (true) {  
    wait (wrt) ;  
  
    //  writing is performed  
  
    signal (wrt) ;  
}
```

# Readers-Writers Problem (Cont.)



- The structure of a reader process

```
while (true) {  
    wait (mutex) ;  
    readcount ++ ;  
    if (readcount == 1)  
        wait (wrt) ;    //( don't allow writers)  
    signal (mutex); //(allow other readers to come)
```

// reading is performed

```
    wait (mutex) ;    //(one by one readers leave the C.S.)  
    readcount - - ;  
    if (readcount == 0)  
    { signal (wrt) ; } // last reader will do this  
        signal (mutex) ; // mutex=1 (for readers to exit and writer to  
come)  
    }
```



**writer**

**wait(wrt);**

...

writing is performed

...

**signal(wrt);**

**reader**

first-in

**wait(mutex);**

**readcount++;**

**if (readcount == 1) wait(wrt);**

**signal(mutex);**

...

reading is performed

...

**wait(mutex);**

**readcount--;**


**if (readcount == 0) signal(wrt);**

**signal(mutex);**



update  
readcount

# Dining-Philosophers Problem

- ◆ Five philosophers, either **thinking** or **eating**
- ◆ to eat, two chopsticks are required
- ◆ taking one chopstick at a time
- ◆ Shared data  
**`semaphore chopstick[5];`**  
Initially all values are 1



# Dining-Philosophers Problem



# Dining-Philosophers Problem

**philosopher  $i$**

do {

**wait(chopstick[i]);**

**wait(chopstick[(i+1) % 5]);**

...

eat

...

**signal(chopstick[i]);**

**signal (chopstick[(i+1) % 5]);**

...

think

...

} while(1);

get chopsticks

left

right

free chopsticks

left

right

**deadlock !**

# Dining-Philosophers Problem

## ◆ Possible solutions to the deadlock problem

- Allow at most **four** philosophers to be sitting simultaneously at the table.
- Allow a philosopher to pick up her chopsticks only if **both chopsticks are available** (note that she must pick them up in a critical section).
- Use an asymmetric solution; that is,
  - ◆ odd philosopher: **left first, and then right**
  - ◆ an even philosopher: **right first, and then left**

## ◆ Besides deadlock, any satisfactory solution to the DPP problem must avoid the problem of **starvation**.

# Monitors

- ❑ A way to encapsulate the Critical Section by making class around the critical section and allowing only one process to be active in that class at one time.
- ❑ The monitor type is a high-level synchronization construct.
- ❑ Only one process may be active within the monitor at a time
  - ❑ Name of Monitor
  - ❑ Initialization Code Section
  - ❑ Procedure to request the Critical Data
  - ❑ Procedure to release the Critical Data

```
monitor monitor name
{
    // shared variable declarations

    procedure P1 ( . . . ) {
        . . .
    }

    procedure P2 ( . . . ) {
        . . .
    }


    .
    .
    .
    procedure Pn ( . . . ) {
        . . .
    }

    initialization code ( . . . ) {
        . . .
    }
}
```

**Figure 6.16** Syntax of a monitor.

# Schematic view of a Monitor

Only one process at a time can be in monitor





# Monitors

- ❑ The **monitor** construct is not sufficiently powerful for modeling some synchronization schemes.
- ❑ So, we need to define additional synchronization mechanisms.
- ❑ These mechanisms are provided by the **condition construct**.
- ❑ A programmer can define one or more variables of type **condition**:

**condition x, y;**

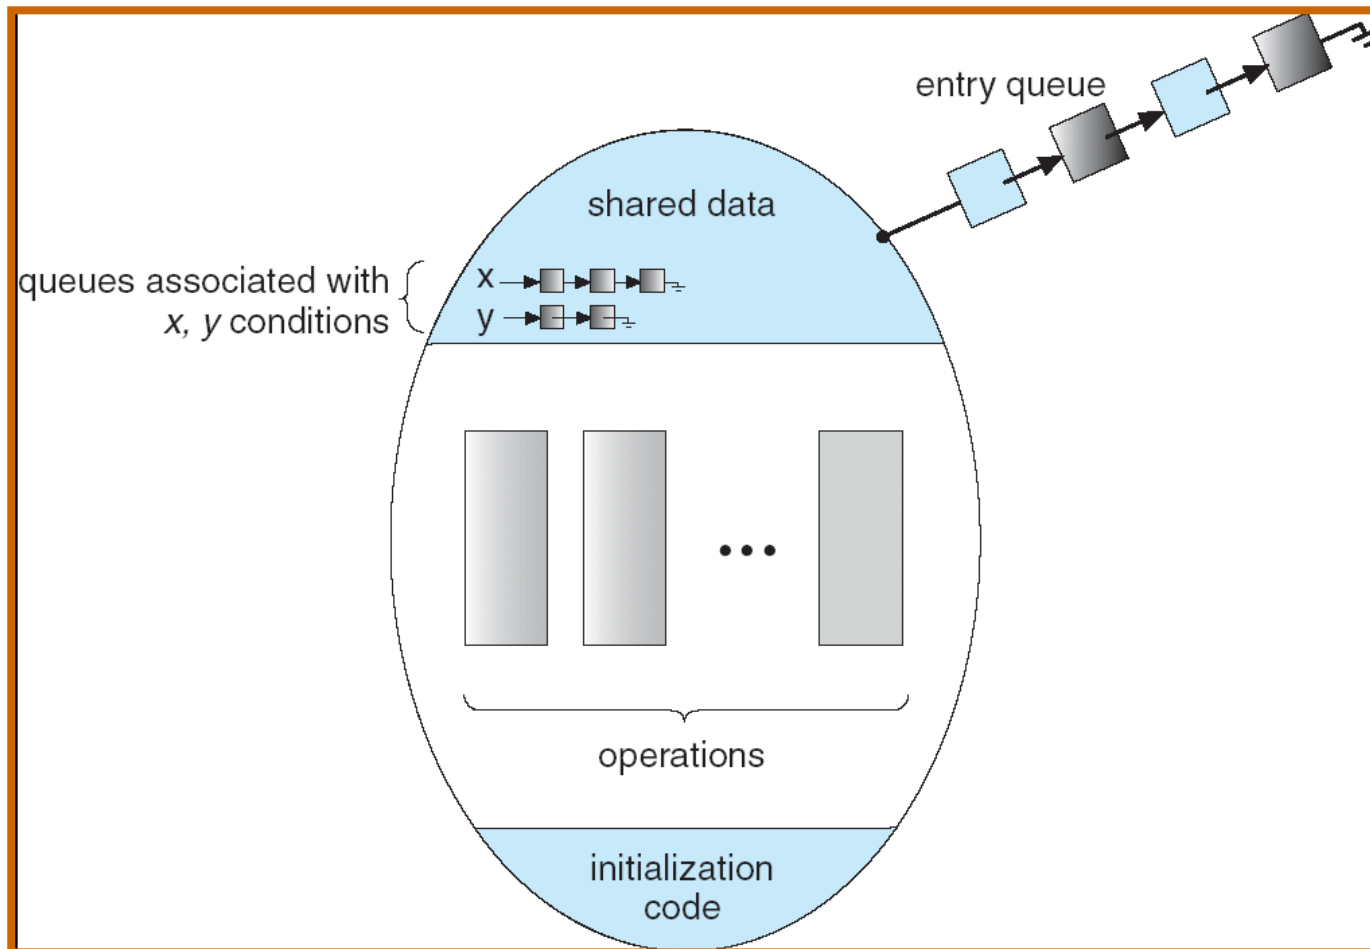
- ❑ The only operations that can be invoked on a **condition variable** are **wait ()** and **signal()**.

# Condition Variables

- Two operations on a condition variable:
  - **x.wait ()** – a process that invokes the operation is suspended until another process invokes **x.signal()**
  - **x.signal ()** – resumes one of processes (if any) that invoked **x.wait ()**

There could be different conditions for which a process could be waiting

# Monitor with Condition Variables



# Solution to Dining Philosophers (cont)

- The distribution of the chopsticks is controlled by the **monitor**

DiningPhilosophers

- Each philosopher ' i ' invokes the operations `pickup()` and `putdown()` in the following sequence:

`dp.pickup (i)`

EAT

`dp.putdown (i)`

# Solution to Dining Philosophers

monitor DP

```
{
    enum { THINKING; HUNGRY, EATING} state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            self [i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```

# Solution to Dining Philosophers (cont)

```
void test (int i) {  
    if ( (state[(i + 4) % 5] != EATING) && (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) )  
{  
        state[i] = EATING ;  
        self[i].signal() ;  
    }  
}  
  
initialization_code() {  
    for (int i = 1; i <= 5; i++)  
        state[i] = THINKING;  
}  
}
```

