

Chapter : Process Synchronization

Background

- ❑ **Co-Operating Process:** that can affect or be affected by other processes executing in system
- ❑ Concurrent access to shared data may result in data inconsistency
- ❑ **Process Synchronization:** Ensures coordination among processes and maintains Data Consistency

❑ Process P1

1. $X=5$

2. $X=5+2$

Process P2

1. $\text{read}(x);$

2. $x=x+5;$

3. $\text{Printf}(x);$



Producer Consumer Problem

There can be two situations:

- 1. Producer Produces Items at Fastest Rate Than Consumer Consumes**
- 2. Producer Produces Items at Lowest Rate Than Consumer Consumes**



Producer Consumer Problem

**Producer Produces Items at Fastest Rate Than
Consumer Consumes:**

If Producer produces items at fastest rate than
Consumer consumes Then Some items will be
lost

Eg. Computer → Producer


Printer → Consumer

Producer Consumer Problem

Solution:

To avoid mismatch of items Produced or Consumed →
Take Buffer

Idea is: Instead of sending items from Producer to
Consumer directly → Store items into buffer



Producer Consumer Problem

Buffer Can be:

1. Unbounded Buffer:

1. No buffer size limit
2. Any no. of items can be stored
3. Producer can produce on any rate, there will always be space in buffer

2. Bounded Buffer:

1. Limited buffer size



Producer Consumer Problem

Bounded Buffer:

If rate of Production $>$ rate of Consumption:

Some items will be unconsumed in buffer

If rate of Production $<$ rate of Consumption:

At some time buffer will be empty

Producer

```
while (true) {
```

```
    /* produce an item and put in  
    nextProduced */
```

```
    while (count == BUFFER_SIZE)  
        ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;
```

```
}
```


Consumer

```
while (true) {  
    while (count == 0)          // buffer empty  
    {  
        ; // do nothing  
    }  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count- -;  
  
    /* consume the item in nextConsumed  
}
```

Race Condition

- When multiple processes access and manipulate the same data at the same time, they may enter into a race condition.
- **Race Condition: When output of the process is dependent on the sequence of other processes.**
- Race Condition occurs when processes share same data

□ Process P1

1. reads $i=10$
2. $i=i+1 =11$

3. Stores $i=11$ in memory

Process P2

1. P2 reads $i=11$ from memory
2. $i=i+1 = 12$
3. Stores 12 in memory

Critical Section Problem

- Section of code or set of operations, in which process may be changing shared variables, updating common data.
- **A process is in the critical section if it executes code that manipulate shared data and resources.**
- Each process should seek permission to enter its critical section → **Entry Section**
- **Exit Section**
- **Remainder section:** Contains remaining code

Structure of a process

Repeat

{

Locks are set here

// Entry Section

Critical Section (a section of code
where processes work with shared
data)

Critical Section

Locks are released here

// Exit Section

Remainder Section

} until false.

Solution to: Critical Section

1. Mutual Exclusion:

It states that **if one process is executing in its critical section**, then no other process can execute in its critical section.

2. Bounded Wait:

It states that if a process makes a request to enter its critical section and before that request is granted, **there is a limit on number of times other processes are allowed to enter that critical section**.

3. Progress:

It states that **process cannot stop other process from entering their critical sections, if it is not executing in its CS**.

Peterson's Solution

- Two process solution (Software-based)
- The two processes share two variables:
 - int **turn**;
 - Boolean **flag[2]**
- The variable **turn** indicates whose turn it is to enter the critical section.
- The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i]=true** implies that process P_i is ready!

- Only **2 processes**,
- General structure of process P_i (other process P_j)
do {
 entry section
 critical section
 exit section
 reminder section
} **while (1);**

Algorithm for Process P_i

do {

flag[i] = TRUE;

turn = j;

while (flag[j] && turn == j);

CRITICAL SECTION

flag[i] = FALSE;

REMAINDER SECTION

do {

acquire lock

critical section

release lock

remainder section

} while (TRUE);

Critical Section Problem solution

- Now prove that this solution is correct.
We need to show that:
 - Mutual exclusion is preserved.
 - The progress requirement is satisfied.
 - The bounded-waiting requirement is met.

Synchronization Hardware

Hardware Solution to C.S.

- ❑ Many systems provide hardware support for critical section code
- ❑ Uni-processors – could disable interrupts
 - ❑ Currently running code would execute without preemption
- ❑ Modern machines provide special atomic hardware instructions
 - ▶ **Atomic = non-interruptable**
 - ❑ Either test memory word and set value
 - ❑ Or swap contents of two memory words

1. Interrupt Disabling

1. Process leaves control of CPU when it is interrupted.

2. Solution is:

- 1. To have each process disable all interrupts just after entering to the critical section.**
- 2. Re-enable interrupts after leaving critical section**

□ Interrupt Disabling

Repeat

Disable interrupts

C.S

Enable interrupts

Remainder section

Synchronization Hardware

- Having the support of some simple hardware instructions, the CS problem can be solved very easily and efficiently.
- The CS problem occurs because the modification of a shared variable of a process may be interrupted.
- Two common hardware instructions that execute **atomically**
 - **Test-and-Set**
 - **Swap**

Synchronization Hardware

- Test and modify the content of a word atomically.

```
boolean TestAndSet(boolean &target) {  
    boolean rv = target;  
    target = true;  
  
    return rv;  
}
```

Mutual Exclusion with Test-and-Set

- Shared data:
boolean lock = false;
- Process P_i
do {
 while (TestAndSet(lock)) ;
 critical section
 lock = false;
 remainder section
}

Synchronization Hardware

- Atomically swap two variables.

```
void Swap(boolean &a, boolean  
&b) {  
    boolean temp = a;  
    a = b;  
    b = temp;  
}
```

Mutual Exclusion with Swap



- Shared data (initialized to **false**):
boolean lock; /***global** variable
- Process P_i
do {
 key = true;
 while (key == true)
 Swap(lock, key);
 critical section
 lock = false;
 remainder section
}

Semaphore

- Synchronization tool that maintains concurrency using variables
- Semaphore S is a integer variable which can take positive values including 0. It is accessed from 2 operations only.

Operations On Semaphore:

- `wait()` and `signal()`

1. Wait Operation is also known as P() which means to test
2. Signal Operation is also known as V() which means to increment

- Entry to C.S. is controlled by `wait()`
- Exit from C.S. is signaled by `signal()`

- Can only be accessed via two indivisible (atomic) operations and $S=1$
- For critical Section problem semaphore value is always 1
- $P(S)$ and $V(S)$:

```
wait (S) {  
    while (S <= 0)  
    do skip ; // no-op}  
    S- -;
```

C.S

```
signal (S) {  
    S++;    }
```

Semaphore as General Synchronization Tool

Types of Semaphores:

- **Counting** semaphore – when integer value can be any non-negative value
- **Binary** semaphore – integer value can range only between 0 and 1
 - Also known as **mutex locks**

Semaphore and Busy Waiting

Disadvantage of Semaphore: Busy Waiting

- ❑ When a process is in C.S. and any other process that wants to enter C.S. loops continuously in entry section.
- ❑ Wastes CPU cycles
- ❑ Semaphore that implements busy waiting is known as: **Spin Lock**

Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated **waiting queue**. Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list



Semaphore Implementation with no Busy waiting

Instead of waiting, a process blocks itself.

- Two operations:
 - **block** – place the process invoking the operation on the waiting queue.
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue.

Semaphore Implementation with no Busy waiting

□ Implementation of wait:

$S=1$

```
wait (S){  
    value--;  
    if (value < 0) {  
        add process P to waiting queue  
        block();  
    }  
C.S  
}
```

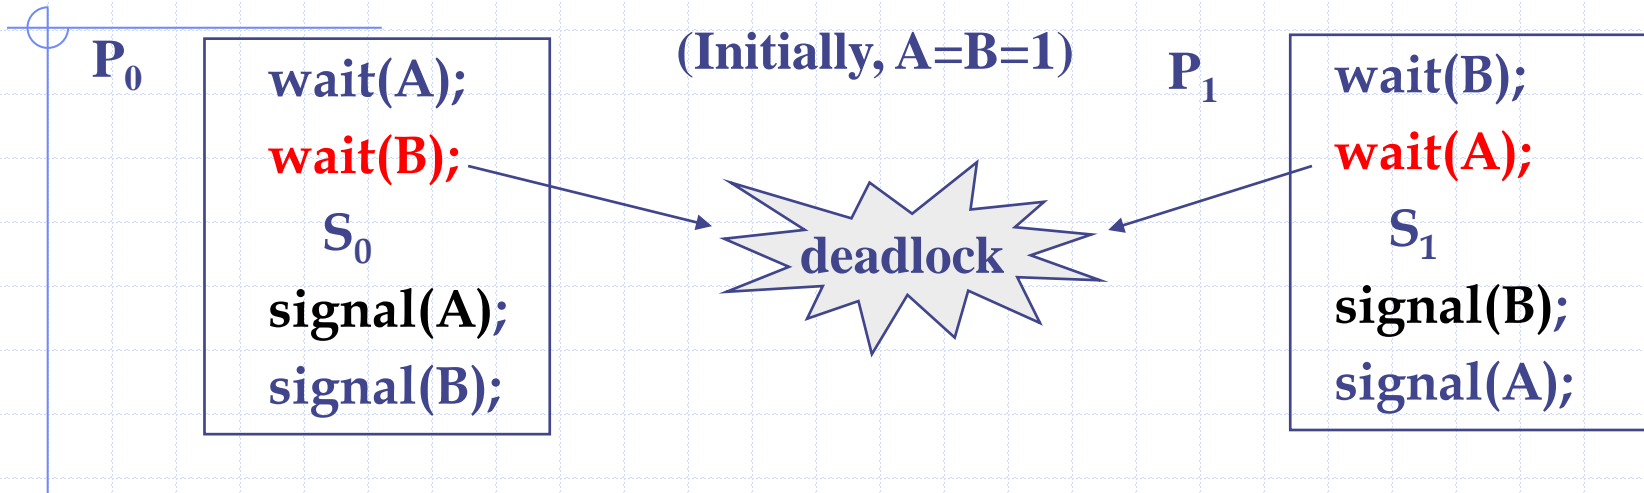
Semaphore Implementation with no Busy waiting

- Implementation of signal:

```
Signal (S){  
    value++;          //“value” has -ve value i.e -1 here  
    if (value <= 0) {  
        remove a process P from the waiting queue  
        wakeup(P); }  
}
```

Deadlocks and Starvation

- ◆ The using of semaphores may cause **deadlocks**



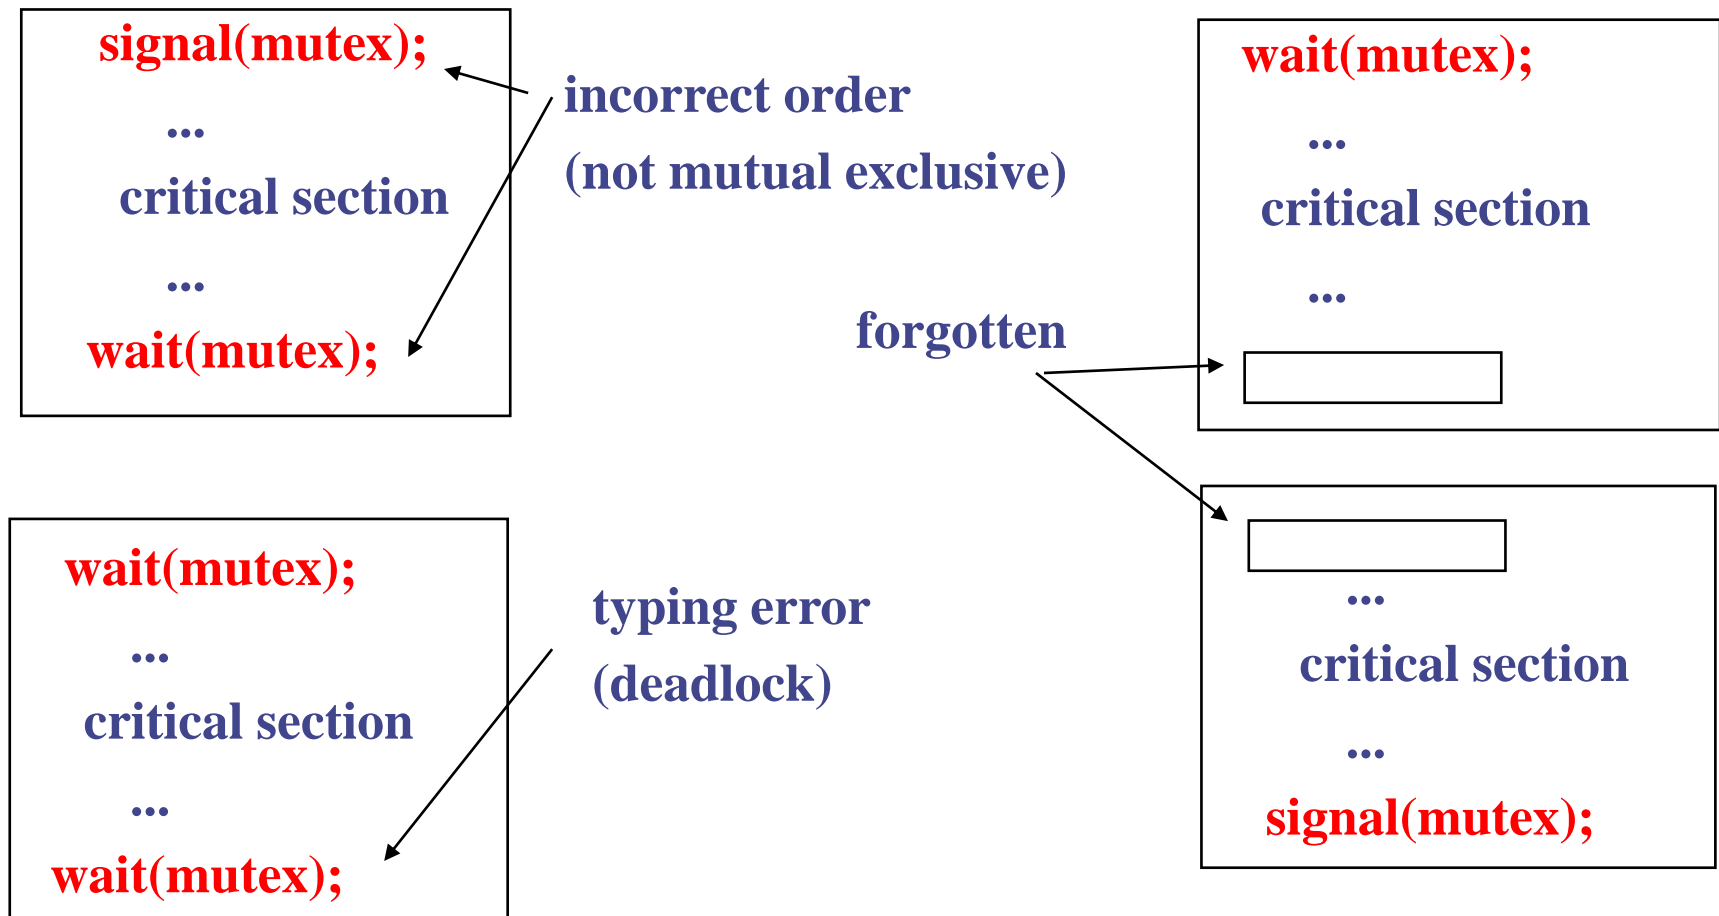
- ◆ **Starvation** – **indefinite blocking**. A process may never be removed from the semaphore queue in which it is suspended

For example: waiting queues are implemented in LIFO order.

Drawbacks of Semaphores



- Semaphores provide a convenient and effective mechanism for process synchronization.
- However, **incorrect use** may result in timing errors.



Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - Readers – only read the data set; they do **not** perform any updates
 - Writers – can both read and write.
- Problem – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.
- Shared Data
 - Data set
 - For Readers: Semaphore **mutex** initialized to 1.
 - For Writers: Semaphore **wrt** initialized to 1.
 - Integer **readcount** initialized to 0.

Readers-Writers Problem (Cont.)

- The structure of a writer process

```
while (true) {  
    wait (wrt) ;  
  
    //  writing is performed  
  
    signal (wrt) ;  
}
```

Readers-Writers Problem (Cont.)



- The structure of a reader process

```
while (true) {  
    wait (mutex) ;  
    readcount ++ ;  
    if (readcount == 1)  
        wait (wrt) ;    //( don't allow writers)  
    signal (mutex); //(allow other readers to come)
```

 // reading is performed

```
    wait (mutex) ;    //(one by one readers leave the C.S.)  
    readcount - - ;  
    if (readcount == 0)  
    { signal (wrt) ; } // last reader will do this  
    signal (mutex) ; // mutex=1 (for readers to exit and writer to  
come)  
}
```


writer

wait(wrt);

...

writing is performed

...

signal(wrt);

reader

first-in

wait(mutex);

readcount++;

if (readcount == 1) wait(wrt);

signal(mutex);

...

reading is performed

...

wait(mutex);

readcount--;


if (readcount == 0) signal(wrt);

signal(mutex);



update
readcount

Dining-Philosophers Problem

- ◆ Five philosophers, either **thinking** or **eating**
- ◆ to eat, two chopsticks are required
- ◆ taking one chopstick at a time
- ◆ Shared data
`semaphore chopstick[5];`
Initially all values are 1



Dining-Philosophers Problem



Dining-Philosophers Problem

philosopher i

do {

wait(chopstick[i]);

wait(chopstick[(i+1) % 5]);

...

eat

...

signal(chopstick[i]);

signal (chopstick[(i+1) % 5]);

...

think

...

} while(1);

get chopsticks

left

right

free chopsticks

left

right

deadlock !

Dining-Philosophers Problem

◆ Possible solutions to the deadlock problem

- Allow at most **four** philosophers to be sitting simultaneously at the table.
- Allow a philosopher to pick up her chopsticks only if **both chopsticks are available** (note that she must pick them up in a critical section).
- Use an asymmetric solution; that is,
 - ◆ odd philosopher: **left first, and then right**
 - ◆ an even philosopher: **right first, and then left**

◆ Besides deadlock, any satisfactory solution to the DPP problem must avoid the problem of **starvation**.

Monitors

- ❑ A way to encapsulate the Critical Section by making class around the critical section and allowing only one process to be active in that class at one time.
- ❑ The monitor type is a high-level synchronization construct.
- ❑ Only one process may be active within the monitor at a time
 - ❑ Name of Monitor
 - ❑ Initialization Code Section
 - ❑ Procedure to request the Critical Data
 - ❑ Procedure to release the Critical Data

```
monitor monitor name
{
    // shared variable declarations

    procedure P1 ( . . . ) {
        . . .
    }

    procedure P2 ( . . . ) {
        . . .
    }


    .
    .
    .
    procedure Pn ( . . . ) {
        . . .
    }

    initialization code ( . . . ) {
        . . .
    }
}
```

Figure 6.16 Syntax of a monitor.

Schematic view of a Monitor

Only one process at a time can be in monitor



Monitors

- ❑ The **monitor** construct is not sufficiently powerful for modeling some synchronization schemes.
- ❑ So, we need to define additional synchronization mechanisms.
- ❑ These mechanisms are provided by the **condition construct**.
- ❑ A programmer can define one or more variables of type **condition**:

condition x, y;

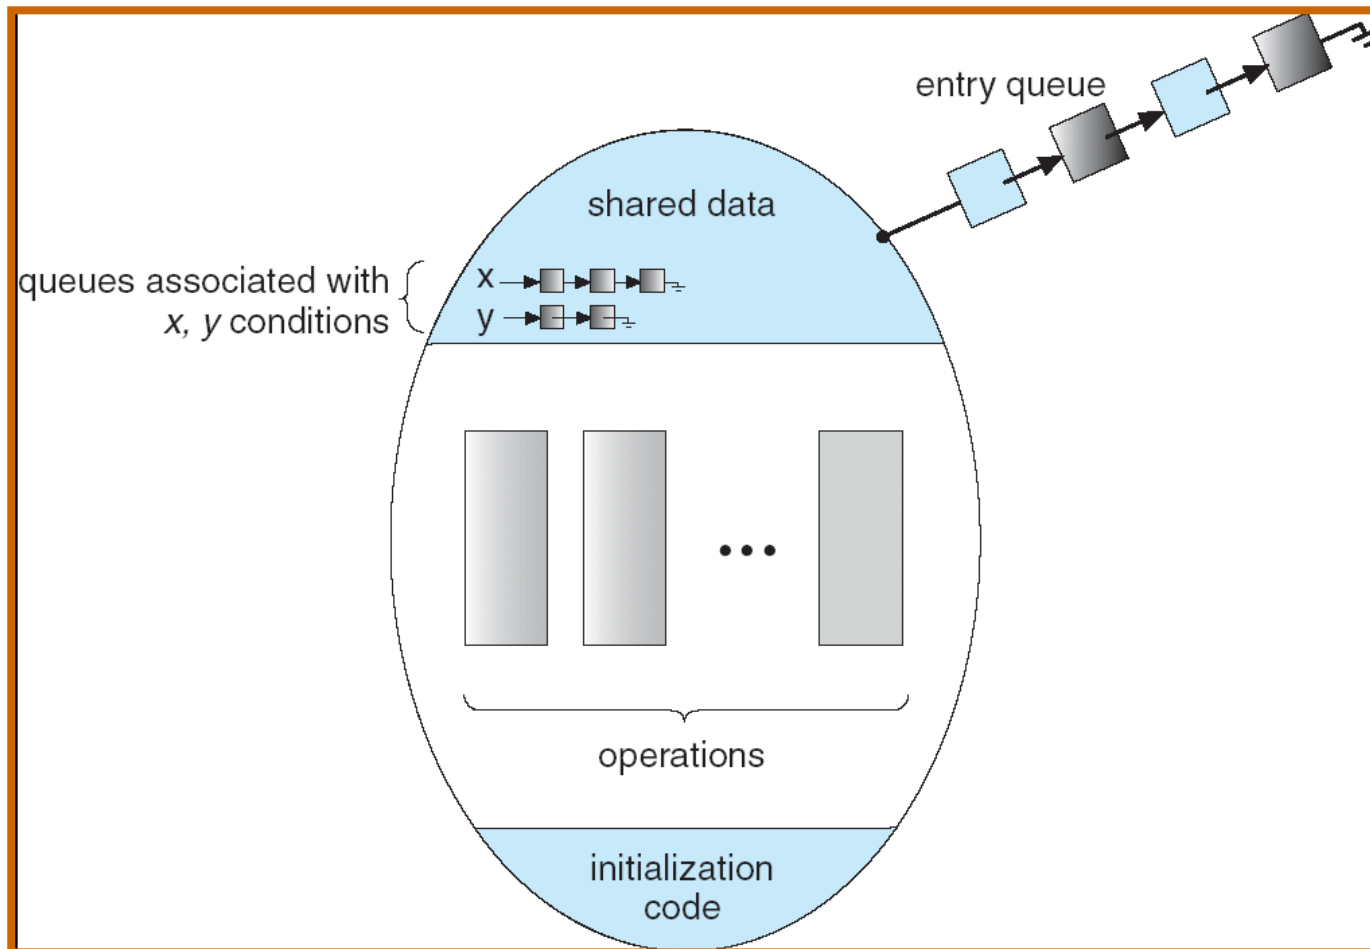
- ❑ The only operations that can be invoked on a **condition variable** are **wait ()** and **signal()**.

Condition Variables

- Two operations on a condition variable:
 - **x.wait ()** – a process that invokes the operation is suspended until another process invokes **x.signal()**
 - **x.signal ()** – resumes one of processes (if any) that invoked **x.wait ()**

There could be different conditions for which a process could be waiting

Monitor with Condition Variables



Solution to Dining Philosophers (cont)

- The distribution of the chopsticks is controlled by the **monitor**

DiningPhilosophers

- Each philosopher ' i ' invokes the operations `pickup()` and `putdown()` in the following sequence:

`dp.pickup (i)`

EAT

`dp.putdown (i)`

Solution to Dining Philosophers

monitor DP

```
{
    enum { THINKING; HUNGRY, EATING) state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            self [i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```



Solution to Dining Philosophers (cont)

```
void test (int i) {  
    if ( (state[(i + 4) % 5] != EATING) && (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) )  
{  
        state[i] = EATING ;  
        self[i].signal() ;  
    }  
}  
  
initialization_code() {  
    for (int i = 1; i <= 5; i++)  
        state[i] = THINKING;  
}  
}
```

